FuxiPulse

Checkpoint 1 Decision Gate
Mentor Presentation Script & Q&A

ENT207TC Digital Startup Lab · Week 4
MVP v0.1 Clickable Prototype
Product: AI decision-support for homestay weak-season operations

How to use this document
Speak the italic quoted lines aloud. Bracketed [placeholders] should be replaced with your real interview quotes, peer feedback, or research sources before the meeting. Target length: **8–10 minutes** + Q&A.

0. Opening (~30 seconds)

“Hello — we are the FuxiPulse team. Today we are not claiming a finished product. We are asking for a clear decision: Is this problem worth continuing? Is the solution concept clear? Is the prototype realistic enough to start building? I will cover four parts — problem, validation evidence, solution concept, and a live clickable demo — and end with an explicit continue / revise / pivot ask.”

1. The Problem (~1.5 minutes)

Who experiences it

“The primary users are independent and small-chain homestay operators and property managers in tourist destinations. They rely heavily on OTAs such as Ctrip, Meituan, and Fliggy; they usually have little or no dedicated revenue-ops staff; and they must make pricing and promotion decisions under time pressure in the weak or shoulder season. Our demo property is Cloud Mist Homestay — a 12-room property in Tengchong, Yunnan — standing in for a typical OTA-dependent independent homestay, not a large hotel chain with a full revenue team.”

What exactly happens

“In weak or shoulder season, occupancy and revenue drop sharply. Booking pace often softens early, but owners usually notice too late — after empty rooms and slipping channel ranks have already compounded. The core gap is not a lack of data. Operators already have OTA calendars, guest messages, and competitor listings. What they lack is a timely, connected way to: see weak-season risk early; diagnose why demand is falling; decide what to do on pricing and promotions; and act across OTA channels before losses widen.”

Why it matters

“In weak season, small delays are expensive. When forward bookings fall and weekday demand softens, competitors can cut rates, listing ranks slip, and conversion drops across channels. Without early diagnosis and coordinated OTA action, revenue shortfalls stack up over days and weeks. This is not “another pretty dashboard” — it is a real operational timing and decision problem.”

2. Problem Validation (~2 minutes)

This is the heart of the decision gate: evidence that the problem is not imagined.

Evidence A — Industry / structural observation

“Independent homestays depend on OTAs for discovery and bookings. Seasonality in tourist destinations is well known. Operators juggle multiple apps for calendars, prices, and ranks, without a revenue team to synthesise signals into one decision path. That friction is structural, not hypothetical.”

Evidence B — Tool / workflow gap

“Existing channel back-ends show “rooms are empty today,” but they rarely answer in one place: Is this weekday softness, a price gap, or conversion drop? Should we push a package or cut rates? How do we sync the decision across channels? The pain is not “need more charts” — it is a missing detect-to-action loop.”

Evidence C — Conversations / peer feedback / research

“Replace with your real evidence before presenting: [peer Green/Yellow/Red outcome]; [quotes from X operators or hospitality contacts — e.g. “we adjust prices by gut feel in weak season,” “we hesitate to cut rates across every channel,” “weekdays are harder than weekends”]; [1–2 secondary sources on Yunnan/Tengchong seasonality or small-lodging revenue management].”

Honest boundary (say this — it builds trust)

“Numbers in the prototype — for example occupancy around 43%, risk score 78/100, roughly ¥18.6k shortfall over two weeks if no action — are illustrative mock data for Cloud Mist Homestay. They frame the decision problem; they are not finished field validation. Our next step is to strengthen interviews, seasonality sources, and to separate validated claims from demo metrics. Today we ask whether the problem direction is sound enough to continue validating and building — not whether every KPI is already proven.”

3. Solution Concept (~1 minute)

“One-line concept: Our solution will help independent and small-chain homestay operators solve late, fragmented weak-season decision-making by turning scattered OTA and demand signals into an AI-assisted detect → diagnose → research → approve → act → learn operations loop.”

“FuxiPulse is an AI operations decision-support dashboard. It does not replace the operator. It compresses scattered signals so a small team can detect risk, understand causes, choose a tactic, push approved actions to OTAs, and learn from results. Design principle: no automatic price change without human confirmation. AI proposes; humans decide — because owners carry the margin risk.”

4. Clickable Prototype Demo — MVP v0.1 (~3–4 minutes)

Core journey to click live: **Dashboard** → **Weak Season** → **Research** → **Strategy (Approve)** → **OTA** → **Feedback**

① Operations Dashboard

“This is the entry point. One glance answers: Are we in trouble? You see property pulse — occupancy, ADR, trends — a composite risk score, and agent pipeline status. Operators should not need five OTA apps just to feel the risk.”

② Weak Season Analysis (Predict & Diagnose)

“Knowing occupancy is low is not enough. Here we show a near-term ~14-day demand view and root-cause diagnosis — for example weekday tourism softness, price gap versus competitors, or falling OTA conversion. The goal is a targeted response: do not cut weekend rates when the real gap is midweek demand.”

③ Research Agent

“Small teams lack time for desk research in a weak-season crunch. This screen surfaces sources, a short reasoning chain, and findings from tourism signals, weather, events, OTA market context, and hospitality playbooks — so recommendations are explainable, not a black box.”

④ Strategy Center + Human Approval

“This is the decision heart. The system ranks tactics — for example a weekday autumn package, a selective weekday rate adjustment, or content/ranking boost — with expected impact and risk. The operator can Approve, Modify, or Reject. Control stays with the human.”

⑤ OTA Operations

“Diagnosis without execution does not stop losses. After approval, actions move to the channel layer: monitor share, booking pace, and listing rank, and sync approved pricing or package actions across OTAs.”

⑥ Knowledge Base + Performance Feedback (brief)

“Knowledge holds reusable playbooks across pricing, product, distribution, and demand. Feedback compares before/after metrics and feeds outcomes into the next cycle — so FuxiPulse becomes an evolving ops system, not a one-shot report. MVP scope note: v0.1 prioritises this core path and seven screens; full live OTA APIs, complex permissions, and multi-property scale are later iterations — not what this checkpoint must prove.”

5. Close — Ask for the Decision (~1 minute)

“To summarise three things: First, the problem is meaningful — OTA-dependent small homestays react late in weak season, and losses compound. Second, the solution is clear and buildable — an AI-assisted ops loop with human approval, scoped for this semester. Third, the prototype is clickable — from risk detection through approval to channel action, with features chosen for a reason.”

“The decision we request today: Should this team continue developing FuxiPulse? If Continue — we will strengthen real-user validation and push the prototype toward buildable core features. If Revise — please tell us which user, scenario, or feature to narrow. If Pivot — we would rather discover that in Week 4 than in Week 9. Thank you — we welcome your questions.”

Q&A — Anticipated Mentor Questions

Use these as short spoken answers. Keep each reply under ~45 seconds unless the mentor digs deeper.

Q1. How is this different from a hotel RMS / revenue management system?

A: We target independent and small-chain homestays without a dedicated revenue team. The product emphasises explainable research, human approval, and an OTA execution loop — not fully automated enterprise pricing for large hotels.

Q2. Is your validation strong enough to continue?

A: Directional evidence exists: structural OTA dependence, the detect-to-action gap, and [peer / interview / research evidence you insert]. Quantitative KPIs in the UI are still demo metrics. We treat that honestly and plan interviews plus seasonality sources next. Continuity should be judged on problem quality and buildable scope, not finished field stats.

Q3. Why insist on human-in-the-loop?

A: Pricing and promotions directly affect margin and channel reputation. Owners carry the risk, so they need final control. It also improves trust and makes the product credible for real operators and for classroom demonstration.

Q4. Can you realistically build this in one semester?

A: Yes for a focused MVP. v0.1 already walks the core journey. Next we prioritise detect → diagnose → strategy approval → feedback. OTA sync can start as mocked or semi-manual, then deepen. Supporting screens stay lightweight until the core loop is solid.

Q5. Who is the first paying user, and what is the wedge?

A: First wedge: independent OTA-dependent homestays in seasonal destinations (demo: Tengchong-style properties). Pain wedge: early weak-season risk + one approved tactic (e.g. weekday package) rather than a full hotel suite on day one.

Q6. What if AI recommendations are wrong?

A: That is why we separate Research (sources + reasoning) from Strategy (ranked options) and require Approve / Modify / Reject. Wrong suggestions should be visible and stoppable. Feedback then improves the next cycle.

Q7. Why seven features — isn’t that too broad for Week 4?

A: Seven screens map one-to-one to one closed loop. The demo path is still one journey. Knowledge and Feedback can be shown briefly; the decision gate is the approve-then-act spine, not feature completeness.

Q8. What would make you pivot?

A: Clear signals that operators do not feel late diagnosis as painful; that they refuse any AI-assisted pricing help; or that the buildable scope collapses into a static report with no action loop. We would rather pivot early than polish a weak project into Week 9.

Q9. Show me the user journey in one sentence.

A: Open Dashboard → see elevated risk → diagnose weekday softness → review research → approve a weekday autumn package → push to OTAs → check whether occupancy and RevPAR moved.

Q10. What do you need from me as Mentor today?

A: A clear gate decision — Continue, Revise, or Pivot — plus the sharpest risk you see in our problem framing, validation gaps, or MVP scope so we know what to fix before we invest the next weeks of build time.

Quick-reference Q&A table

Topic	Short answer
vs hotel RMS	Small OTA-dependent homestays; explainable + human approve + OTA act
Validation level	Direction yes; UI KPIs are mock — interviews/sources next
Human-in-loop	Owners own margin risk; AI proposes, humans decide
Semester fit	Core loop first; OTA can start mocked / semi-manual
Ask today	Continue / Revise / Pivot + sharpest risk to fix

End of document · Replace bracketed validation placeholders before Mentor meeting · Demo URL / local: `npm run dev` → http://127.0.0.1:5173